

EXECUTIVE STRATEGY

SaNDAI 8-Week Seminar Growth Plan

An auditable operating system for at least 20 actual student attendees every week



RECOMMEND

Run the challenge as a weekly attendance funnel with institutional multipliers, dated decision gates, and a command center that turns current performance into the next action. The campaign is not a one-time promotion; it is an eight-week operating cycle owned by Piter Garcia.

[Open the interactive prototype](#) | No sign-in required. Review steps and supporting evidence appear on page 7.

The decision in one minute

Question	Operating answer
Objective	Starting in October 2026: at least 20 actual attendees in each of 8 consecutive weeks; at least 160 total attendances.
Owner	Piter Garcia - accountable for plan, execution, measurement, intervention, and weekly reporting.
Acquisition model	Anchor relationships first; then faculty, departments, career services, clubs, partners, referrals, and approved digital channels.
Registration requirement	Dynamic: ceiling(20 / conservative show rate), then a configurable safety buffer.
Control system	T-7, T-5, T-3, T-1, seminar day, and T+1 gates with explicit thresholds and actions.
Next decision	Before launch: approve seminar claims, schedule, registration flow, outreach boundaries, and downstream offers.
Immediate next action	Piter - within 24 hours of authorization: confirm operating facts and populate the Week 0 partner pipeline.
Next update	A five-line executive report at every T+1 review; an escalation update whenever status becomes Yellow or Red.

MEASURE **Success is weekly, not merely cumulative**

The 160-attendance total is necessary but not sufficient. A week with 30 attendees does not erase a different week with 15. The command center therefore checks both conditions: every weekly floor and the cumulative total.

Week	1	2	3	4	5	6	7	8
Minimum actual attendance	20	20	20	20	20	20	20	20
Minimum cumulative	20	40	60	80	100	120	140	160

WHY: This design makes ownership visible. A shortfall cannot hide inside an eight-week average, and a weak funnel triggers an action before seminar day.

Ownership lifecycle

Take responsibility	Develop plan	Execute	Measure	Adjust	Deliver
Named owner and objective	Partners, funnel, gates, backups	Dated weekly cadence	Actuals, forecast, channels	Triggered intervention	20/week and final report

REACH

Acquire through multipliers, not endless one-to-one outreach

A layered network that can repeat for eight weeks

Four-layer expansion strategy

Layer	Purpose	Examples and operating move
1. Anchor network	Start quickly and learn	Use approved public/relationship channels through Farmingdale State College (FSC), RIT, and the University of Rochester; do not use private rosters.
2. Institutional multipliers	Reach groups at once	Faculty, departments, career services, student organizations, academic programs, and approved UofR partner or teaching-placement contacts distribute an approved message.
3. Repeat partners	Create a reusable referral engine	Give partners a tracked link or source code, a reminder kit, and a short outcome update so they can refer again.
4. Broader reach	Reduce dependence on campus ties	Professional/community organizations, alumni and peer referrals, public events, and approved digital channels.

Concentration control

By Week 3, the goal is for institutional multipliers, partners, and referrals to supply at least half of registrations. After Week 2, no single channel should supply more than 40% of weekly registrations. These are management thresholds to test and revise, not claims about existing performance.

The UofR relationship is especially valuable because it includes existing partner and teaching-placement relationships that can be explored through approved contacts. Each outreach route still requires the appropriate institutional or supervising contact to approve distribution.

TRACK Transparent attendance funnel

Stage	Decision value	Minimum record
Reached / invited	Shows outreach volume and partner activity	Date, source/channel, campaign code; no unnecessary personal data.
Interested	Identifies message-audience fit	Opt-in or tracked interest signal.
Registered	Creates the first attendance forecast	Registration count by source and seminar.
Confirmed	Protects show rate close to the event	Confirmation response or reminder status.
Attended	Primary KPI and weekly pass/fail result	Aggregate attendance and approved check-in evidence.
Converted	Connects free attendance to approved value	Opt-in follow-up and SaNDAI-approved outcome attribution.

Dynamic registration target

Registration floor = ceiling(20 / conservative show rate). Before live evidence, the planning rate is 60%. After results exist, the notebook uses the 80%-confidence lower estimate from completed registrations and attendance, subtracts 5 percentage points, and applies a 40% floor. The safety target is ceiling(registration floor x 1.10). These assumptions remain editable and visible.

Conservative show rate	Registration floor	Meaning
75%	27	Strong show rate; still protect confirmations.
60%	34	Practical planning baseline before campaign-specific evidence.
50%	40	Weak show rate; increase confirmations and investigate friction.

OPERATE

Weekly decision gates turn data into action

Initial thresholds are explicit, auditable, and recalibrated when real evidence arrives

Gate	Metric / initial threshold	Owner / if healthy	If under target
T-7	At least 3 active multiplier partners; expected contribution at least 20	Piter: release approved partner kit and tracked links	Activate backup prospects; add faculty/career-services lane
T-5	Registrations at least 70% of the 10% safety target	Piter: continue best channels; prepare confirmations	Open secondary partner lane; revise weak audience-message match
T-3	Forecast at least 22; P(at least 20) at least 80%; controls complete	Piter: protect show rate; avoid unnecessary volume	Run targeted confirmation; activate high-yield backup channel
T-1	Forecast remains at least 20 after current confirmations	Piter: send logistics reminder and access support	Use approved final reminder/referral wave; escalate gap and fallback
T	At least 20 actual attendees	Piter: deliver; capture approved aggregate outcomes	Record miss honestly; diagnose stage/channel failure without hiding it
T+1	Complete data, retrospective, next-week target and owner	Piter: reallocate toward repeatable channels	Correct source/process gaps; launch recovery experiment

Forecast status

RED: expected attendance below 20 or a supplied control flag is incomplete. GREEN: expected attendance at least 22 and probability of 20 or more at least 80%. YELLOW: otherwise. The demo assumes controls are complete; it does not inspect them. T-7, T-5, and T-1 gates remain operator checks, not automated verification.

Recovery playbook

Observed problem	Likely stage	Immediate controlled response
Reach is low	Partner pipeline	Add multipliers rather than mass individual messaging; activate backup organizations.
Interest is low	Message / audience fit	Test one revised value proposition with a defined audience; preserve the approved core claims.
Registrations are low	Call-to-action friction	Simplify link and instructions; ask partners to resend through the best-performing channel.
Show rate is low	Confirmation / access	Use confirmation sequence, clear logistics, calendar add, accessibility check, and targeted follow-up.
One channel dominates	Concentration risk	Cap marginal effort there and develop the next-best repeatable source.
Week misses 20	Outcome	Report the miss, isolate the funnel failure, run one recovery experiment, and increase the next registration requirement as warranted.

Five-line executive update

ILLUSTRATIVE - SYNTHETIC, NOT LIVE CAMPAIGN DATA

1) Forecast: 17 attendees vs. 20 minimum. 2) Driver: registrations are adequate, but current show-rate protection is weak. 3) Risk: RED; estimated gap is 3 attendees. 4) Action/owner: Piter activates a secondary partner and targeted confirmation sequence today. 5) Next decision: review revised forecast at T-1.

ANALYZE

Measure what changes the next decision

Channel evidence, partner follow-up, and ethical value conversion

Command-center measures

View	Measures	Decision
Weekly outcome	Actual attendance; weekly pass/fail; cumulative total	Did the campaign achieve the non-negotiable outcome?
Funnel health	Reached, interested, registered, confirmed, attended	Where is participation being lost?
Forecast	Show rate, required registrations, forecast gap, risk status	What must happen before the next gate?
Channel contribution	Registrations, attendance, show rate, share, outreach effort	Which channels deserve more or less effort?
Partner pipeline	Status, expected contribution, next action, follow-up, actual attendance	Which multipliers can repeat next week?
Value conversion	Opt-in leads, qualified follow-up, approved conversions, attribution	Does the free seminar create approved downstream value?

Count each participant once per weekly seminar. Repeat attendance across weeks is allowed; assign one channel attribution per attendance. The prototype sums supplied channel totals and does not deduplicate individual check-ins.

Channel allocation rule

Prioritize repeatable attendance, not vanity reach. The command center compares actual attendees per outreach-effort unit, show rate, partner repeatability, and concentration risk. One effort unit equals one 30-minute campaign work block. Effort moves toward channels that reliably produce attendance, while at least one controlled experiment remains active to expand reach.

Monetization without charging for the free seminar

Free seminar → opt-in interest → approved SaNDAI opportunity → attributed conversion

The seminar remains free. Monetization is measured only through SaNDAI-approved downstream offerings, follow-up rules, prices, and attribution windows. Until SaNDAI supplies those terms, the command center demonstrates the conversion architecture with synthetic counts and leaves dollar values blank. It does not invent fees, commissions, credentials, or promises.

Release gate: before any monetization follow-up begins, SaNDAI approves the offer, participant eligibility, terms, follow-up cadence, and attribution window.

Stage	Control	Measure
Opt-in interest	Explicit participant choice; approved language	Count and source/channel
Qualified follow-up	SaNDAI-approved criteria and human review	Qualified rate and follow-up completion
Conversion	Only confirmed approved outcome	Count, rate, and source attribution
Economic value	Configured only after rates/terms are confirmed	Approved revenue/value metric; no fabricated assumptions

Minimum data discipline

Use source codes and aggregate campaign measures wherever possible. If SaNDAI later authorizes learner-level improvement analysis, use explicit notice/consent and a pseudonymous participant ID; keep identity mapping in an approved system outside the notebook. Collect only necessary access, checkpoint, feedback, and support signals, with defined purpose, access, retention, and deletion rules. Do not collect passwords, raw authentication logs, or sensitive traits by default. Separate synthetic demonstration data from any future live campaign dataset.

Method: understand, structure, build, validate, improve

CRISP-DM keeps the work anchored in the business objective and the data before analysis. Iterative SDLC then delivers one bounded improvement at a time: plan, build, test, use, measure, and revise. The data architecture moves from raw campaign events to structured and validated records, curated metrics, analysis, and decision-ready guidance. Authorization remains outside that data flow.

Use AI as bounded support, not an unaccountable operator

Proposed AI support roles; the prototype uses sample responses and deterministic analytics, not autonomous agents

Worker	Useful support	Human gate
Research worker	Finds public organizations, programs, and contact routes; organizes evidence	Piter verifies relevance, source, and outreach permission
Message adaptation worker	Adapts an approved core message to audience/context	Piter approves claims, tone, and recipient before external use
Campaign worker	Maintains queues, reminders, partner next actions, and follow-up suggestions	Piter chooses timing and sends consequential communication
Analytics worker	Calculates funnel health, show rate, dynamic targets, forecast, and weak channels	Piter reviews data quality and selects intervention
Learning worker	Summarizes the weekly retrospective and proposes a bounded experiment	Piter approves the experiment and success threshold
Accountable owner	Piter Garcia	Owns outcome, decisions, reporting, escalation, and SaNDAI commitments

Simplest reliable method first

Not every campaign task requires generative AI. Tracked links, forms, scheduled reminders, formulas, and rule-based alerts may be safer and more reliable. AI is used where it adds value beyond ordinary automation, with evidence and human review.

Operational safeguards

- No scraping of private student data and no use of private rosters without authorization.
- No fabricated SaNDAI benefits, credentials, pricing, guarantees, or commercial terms.
- No autonomous mass messaging or commitments on behalf of SaNDAI.
- No unnecessary student PII; authorized learner profiles are pseudonymous, consented, purpose-limited, and access-controlled.
- Human approval before consequential outreach, monetization claims, or campaign changes.
- Approved core content, version history, and a rollback path for message or workflow changes.
- Synthetic demonstration data clearly labeled and never presented as live results.

What the prototype demonstrates

The website makes the two experiences accessible without setup. The supporting notebooks expose the reusable Python logic, inputs, and validation. The proposed operating workflow extends beyond the functions implemented in this prototype.

Experience and supporting notebook	Useful output	Non-negotiable boundary
Management Command Center	Campaign forecast, gates, channel/partner decisions, ethical conversion, and aggregate or pseudonymous learner-support insights	No live tracking without SaNDAI authorization, notice/consent, approved collection method, and data-governance controls
Student Seminar Companion	Private self-check, prompt practice, credential/resume planning, next-step plan, and optional feedback structure	No management data or required personal disclosure. Web planning inputs are sent for calculation without application-level storage; written reflections stay in the browser. The notebook does not call a live AI service.

OWN

Why SaNDAI can trust me to own the eight weeks

Evidence plus an operating system is stronger than adjectives



Direct answer

SaNDAI should not have to rely on me remembering what comes next. The campaign itself contains the target, owner, dated gates, thresholds, alerts, interventions, and next-update cadence. I will make performance visible, act before a miss becomes final, report shortfalls honestly, and document the decision that follows.

Collaboration model — Accountability is centralized; execution is collaborative. I remain accountable for the eight-week objective, operating cadence, measurement, escalation, decisions, and reporting while coordinating approved work across SaNDAI, faculty, career services, student organizations, departments, partners, and any assigned team members. Each delegated action has an owner, deadline, expected output, and follow-up; I remain responsible for tracking the result and triggering the next action when performance falls short.

Evidence I bring to the challenge

Evidence from my work	How it transfers
MLADIS system design	Keeps source-backed facts separate from generated suggestions and places authorized action outside the recommendation layer.
Quantum/AI research workflows	Uses explicit assumptions, reproducible analysis, documented experiments, validation, and revision when evidence changes.
CRISP-DM and iterative SDLC	Understands the objective and data first, then builds, tests, measures, and improves one bounded feature at a time.
Teaching and lesson development	Designs for participation, clear success criteria, checkpoints, access, intervention, and reflection.
Version-controlled project continuity	Keeps sources, changes, current status, and next work visible so progress does not depend on memory.
Responsible AI practice	Separates facts from assumptions, recommendations from authorized action, and automation from accountability.

Pre-launch validation and first operating cycle

When	Action	Output
Before launch	Confirm seminar schedule/content, approved claims, registration process, accessibility, downstream offerings, attribution, and outreach/privacy boundaries	Approved campaign brief and operating definitions
Week 0	Build partner queue, tracked source codes, message kit, dashboard baseline, reminder cadence, and backups	Ready-to-run campaign with owners and dates
Each week	PLAN -> RECRUIT -> CONFIRM -> EXECUTE -> MEASURE -> LEARN -> ADJUST	Twenty-or-more attendance decision loop
After Week 8	Reconcile weekly and cumulative outcomes, partner/channel learning, conversions, risks, and recommendations	Auditable final report and repeatable playbook

Two connected browser experiences

1. Management Command Center

Review weekly attendance, cumulative pace, channels, and follow-on conversion. The sample contains 183 attendances but misses Week 3, so the full eight-week objective is not achieved. The Week 4 forecast and recovery case shows how I would act on risk; estimated improvement is not a guarantee.

2. Student Career Lab

Work through Reflect, Practice, Check, Build, and Plan. Choose a community, career, or teaching scenario; review a sample AI response; explain a change; and download a takeaway with a realistic weekly plan. The website does not call a live AI service or certify completion.

How to review the prototype

Start with the website; no account, installation, or notebook execution is required. This plan explains the strategy and ownership. The browser experience makes the decisions and student value tangible. The repository and notebooks provide inspectable methods and sample inputs.

[1. Open the Management Command Center](#) — select Week 3 to see the miss, then review the Week 4 recovery case and channel contribution.

[2. Try the Student Career Lab](#) — choose a scenario, complete the five review questions, then build and download a takeaway.

[3. Review guide and supporting notebooks](#) — optional Colab notebooks and a reproducible route into the code.

[4. Sample data and configuration](#) — the illustrative inputs; configuration is linked from the review guide.

[5. Requirement traceability](#) — what the prototype demonstrates and what remains a proposed live workflow.

Access note: the public website contains sample data, not a live student campaign. The private RIT Shared Drive supports local and Colab testing; reviewers do not need access to it. No enrollment, outreach, monetization, or student tracking is triggered by visiting the website.

The weekly executive report stays concise

Result/forecast ↔ main driver ↔ current risk ↔ action and owner ↔ next decision gate/update.

This is how I will translate strong analysis into fast executive execution while preserving the evidence discipline SaNDAI expects.

Demonstration boundary: This case is hypothetical. All campaign and learner records are synthetic. Live implementation would begin only after SaNDAI approves the operating facts, outreach boundaries, conversion terms, collection method, consent language, access controls, retention, and student-support workflow.